

Assignment 6: Apply NB

1. Apply Multinomial NB on these feature sets

- **Set 1**: categorical, numerical features + preprocessed_eassay (BOW)
- **Set 2**: categorical, numerical features + preprocessed_eassay (TFIDF)


2. The hyper paramter tuning(find best alpha:smoothing parameter)

- Find the best hyper parameter which will give the maximum AUC (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/receiver-operating-characteristic-curve-roc-curve-and-auc-1/>) value
- find the best hyper paramter using k-fold cross validation(use GridsearchCV or RandomsearchCV)/simple cross validation data (write for loop to iterate over hyper parameter values)


•

3. Representation of results

- You need to plot the performance of model both on train data and cross validation data for each hyper parameter, like shown in the figure



- Once after you found the best hyper parameter, you need to train your model with it, and find the AUC on test data and plot the ROC curve on both train and test.



- Along with plotting ROC curve, you need to print the [confusion matrix](https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/) (<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/confusion-matrix-tpr-fpr-fnr-tnr-1/>) with predicted and original labels of test data points

	Predicted: NO	Predicted: YES
Actual: NO	TN = ??	FP = ??
Actual: YES	FN = ??	TP = ??

- fine the top 20 features from either from feature **Set 1** or feature **Set 2** using absolute values of `feature_log_prob_` parameter of `MultinomialNB` (https://scikit-learn.org/stable/modules/generated/sklearn.naive_bayes.MultinomialNB.html) and print their corresponding feature names
- You need to summarize the results at the end of the notebook, summarize it in the table format

Vectorizer	Model	Hyper parameter	AUC
BOW	Brute	7	0.78
TFIDF	Brute	12	0.79
W2V	Brute	10	0.78
TFIDFW2V	Brute	6	0.78

2. Naive Bayes

1.1 Loading Data

```
In [149]: 1 %matplotlib inline
2 import warnings
3 warnings.filterwarnings("ignore")
4
5 import pandas as pd
6 import numpy as np
7 import nltk
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 from sklearn.feature_extraction.text import TfidfVectorizer
11 from sklearn.feature_extraction.text import CountVectorizer
12 from sklearn.metrics import confusion_matrix
13 from sklearn import metrics
14 from sklearn.metrics import roc_curve, auc
15
16 import re
17 # Tutorial about Python regular expressions: https://pymotw.com/2/re/
18
19 import pickle
20 from tqdm import tqdm
21 import os
22 from collections import Counter
```

```
In [150]: 1 import pandas as pd
2 # data = pandas.read_csv('preprocessed_data.csv')
3 data = pd.read_csv('preprocessed_data.csv', nrows=50000)
4
```

```
In [151]: 1 data.head(1)
```

```
Out[151]:
```

	school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_posted_projects
0	ca	mrs	grades_prek_2	1

1.2 Splitting data into Train and cross validation(or test): Stratified Sampling

```
In [152]: 1 # please write all the code with proper documentation, and proper titles for
2 # go through documentations and blogs before you start coding
3 # first figure out what to do, and then think about how to do.
4 # reading and understanding error messages will be very much helpfull in debugging
5 # when you plot any graph make sure you use
6     # a. Title, that describes your plot, this will be very helpful to the reader
7     # b. Legends if needed
8     # c. X-axis Label
9     # d. Y-axis Label
```

```
In [153]: 1 y = data['project_is_approved'].values
          2 X = data.drop(['project_is_approved'], axis=1)
          3 X.head(1)
```

```
Out[153]:
```

	school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_posted_projects
0	ca	mrs	grades_prek_2	1

```
In [154]: 1 from sklearn.model_selection import train_test_split
          2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, stratify=y)
          3 X_train, X_cv, y_train, y_cv = train_test_split(X_train, y_train, test_size=0.33, stratify=y_train)
          4
          5 print(X_train.shape)
          6 print(X_test.shape)
          7 print("*****")
          8 print(y_train.shape)
          9 print(y_test.shape)
```

```
(22445, 8)
(16500, 8)
*****
(22445,)
(16500,)
```

1.3 Make Data Model Ready: encoding eassy, and project_title

BOW Vectorization

```
In [155]: 1 # please write all the code with proper documentation, and proper titles for
          2 # go through documentations and blogs before you start coding
          3 # first figure out what to do, and then think about how to do.
          4 # reading and understanding error messages will be very much helpfull in debugging
          5 # make sure you featurize train and test data separatly
          6
          7 # when you plot any graph make sure you use
          8     # a. Title, that describes your plot, this will be very helpful to the reader
          9     # b. Legends if needed
         10     # c. X-axis label
         11     # d. Y-axis label
```

BOW Vectorization Eassy

```
In [156]: 1 from sklearn.feature_extraction.text import CountVectorizer
2
3 vect= CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
4 vect.fit(X_train['essay'].values)
5 X_train_essay_bow = vect.transform(X_train['essay'].values)
6 X_cv_essay_bow = vect.transform(X_cv['essay'].values)
7 X_test_essay_bow = vect.transform(X_test['essay'].values)
8 essay_bow_features=vect.get_feature_names()
9
10 print('X_train_essay_bow.shape',X_train_essay_bow.shape, y_train.shape)
11 print('X_test_essay_bow.shape',X_test_essay_bow.shape, y_test.shape)
12 print('X_cv_essay_bow.shape',X_cv_essay_bow.shape, y_cv.shape)
13 print(essay_bow_features[:10])
```

```
X_train_essay_bow.shape (22445, 50000) (22445,)
X_test_essay_bow.shape (16500, 50000) (16500,)
X_cv_essay_bow.shape (11055, 50000) (11055,)
['00', '00 pm', '000', '000 students', '10', '10 000', '10 11', '10 11 year',
'10 15', '10 girls']
```

1.4 Make Data Model Ready: encoding numerical, categorical features

```
In [157]: 1 # please write all the code with proper documentation, and proper titles for
2 # go through documentations and blogs before you start coding
3 # first figure out what to do, and then think about how to do.
4 # reading and understanding error messages will be very much helpfull in debu
5 # make sure you featurize train and test data separatly
6
7 # when you plot any graph make sure you use
8     # a. Title, that describes your plot, this will be very helpful to the re
9     # b. Legends if needed
10    # c. X-axis label
11    # d. Y-axis label
```

school_state -OHE

```
In [158]: 1 vect= CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
2 vect.fit(X_train['school_state'].values)
3 X_train_ohe_school_state= vect.transform(X_train['school_state'].values)
4 X_cv_ohe_school_state= vect.transform(X_cv['school_state'].values)
5 X_test_ohe_school_state= vect.transform(X_test['school_state'].values)
6 school_state_features=vect.get_feature_names()
7
8 print('X_train_ohe_school_state',X_train_ohe_school_state.shape, y_train.shape)
9 print('X_test_ohe_school_state',X_test_ohe_school_state.shape, y_test.shape)
10 print('X_cv_ohe_school_state',X_cv_ohe_school_state.shape, y_cv.shape)
11 print(school_state_features[:10])
```

```
X_train_ohe_school_state (22445, 50) (22445,)
X_test_ohe_school_state (16500, 50) (16500,)
X_cv_ohe_school_state (11055, 50) (11055,)
['ak', 'al', 'ar', 'az', 'ca', 'co', 'ct', 'dc', 'de', 'fl']
```

teacher_prefix -OHE

```
In [159]: 1 vect= CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
2 vect.fit(X_train['teacher_prefix'].values)
3 X_train_ohe_teacher_prefix= vect.transform(X_train['teacher_prefix'].values)
4 X_cv_ohe_teacher_prefix= vect.transform(X_cv['teacher_prefix'].values)
5 X_test_ohe_teacher_prefix= vect.transform(X_test['teacher_prefix'].values)
6 teacher_prefix_features=vect.get_feature_names()
7
8 print('X_train_ohe_teacher_prefix',X_train_ohe_teacher_prefix.shape, y_train.shape)
9 print('X_test_ohe_teacher_prefix',X_test_ohe_teacher_prefix.shape, y_test.shape)
10 print('X_cv_ohe_teacher_prefix',X_cv_ohe_teacher_prefix.shape, y_cv.shape)
11 print(teacher_prefix_features[:10])
```

```
X_train_ohe_teacher_prefix (22445, 4) (22445,)
X_test_ohe_teacher_prefix (16500, 4) (16500,)
X_cv_ohe_teacher_prefix (11055, 4) (11055,)
['mr', 'mrs', 'ms', 'teacher']
```

project_grade_category-OHE

```
In [160]: 1 vect= CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
2 vect.fit(X_train['project_grade_category'].values)
3 X_train_ohe_project_grade_category= vect.transform(X_train['project_grade_cat
4 X_cv_ohe_project_grade_category= vect.transform(X_cv['project_grade_category']
5 X_test_ohe_project_grade_category= vect.transform(X_test['project_grade_categ
6 project_grade_category_features=vect.get_feature_names()
7
8 print('X_train_ohe_project_grade_category',X_train_ohe_project_grade_category
9 print('X_test_ohe_project_grade_category',X_test_ohe_project_grade_category.s
10 print('X_cv_ohe_project_grade_category',X_cv_ohe_project_grade_category.shape
11 print(project_grade_category_features[:10])
12
```

```
X_train_ohe_project_grade_category (22445, 4) (22445,)
X_test_ohe_project_grade_category (16500, 4) (16500,)
X_cv_ohe_project_grade_category (11055, 4) (11055,)
['grades_3_5', 'grades_6_8', 'grades_9_12', 'grades_prek_2']
```

clean_categories-OHE

```
In [161]: 1 vect= CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
2 vect.fit(X_train['clean_categories'].values)
3 X_train_ohe_clean_categories= vect.transform(X_train['clean_categories'].valu
4 X_cv_ohe_clean_categories= vect.transform(X_cv['clean_categories'].values)
5 X_test_ohe_clean_categories= vect.transform(X_test['clean_categories'].values
6 clean_categories_features=vect.get_feature_names()
7
8 print('X_train_ohe_clean_categories',X_train_ohe_clean_categories.shape, y_tr
9 print('X_test_ohe_clean_categories',X_test_ohe_clean_categories.shape, y_test
10 print('X_cv_ohe_clean_categories',X_cv_ohe_clean_categories.shape, y_cv.shape
11 print(clean_categories_features[:10])
12
```

```
X_train_ohe_clean_categories (22445, 36) (22445,)
X_test_ohe_clean_categories (16500, 36) (16500,)
X_cv_ohe_clean_categories (11055, 36) (11055,)
['appliedlearning', 'appliedlearning health_sports', 'appliedlearning history_c
ivics', 'appliedlearning literacy_language', 'appliedlearning math_science', 'a
ppliedlearning music_arts', 'appliedlearning specialneeds', 'health_sports', 'h
ealth_sports appliedlearning', 'health_sports literacy_language']
```

clean_subcategories-OHE

```
In [162]: 1 vect= CountVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
2 vect.fit(X_train['clean_subcategories'].values)
3 X_train_ohe_clean_subcategories= vect.transform(X_train['clean_subcategories']
4 X_cv_ohe_clean_subcategories= vect.transform(X_cv['clean_subcategories'].valu
5 X_test_ohe_clean_subcategories= vect.transform(X_test['clean_subcategories']).
6 clean_subcategories_features=vect.get_feature_names()
7
8 print('X_train_ohe_clean_subcategories',X_train_ohe_clean_subcategories.shape
9 print('X_test_ohe_clean_subcategories',X_test_ohe_clean_subcategories.shape,
10 print('X_cv_ohe_clean_subcategories',X_cv_ohe_clean_subcategories.shape, y_cv
11 print(clean_subcategories_features[:10])
```

```
X_train_ohe_clean_subcategories (22445, 147) (22445,)
X_test_ohe_clean_subcategories (16500, 147) (16500,)
X_cv_ohe_clean_subcategories (11055, 147) (11055,)
['appliedsciences', 'appliedsciences charactereducation', 'appliedsciences coll
ege_careerprep', 'appliedsciences earlydevelopment', 'appliedsciences environme
ntalscience', 'appliedsciences esl', 'appliedsciences extracurricular', 'applie
dsciences health_lifescience', 'appliedsciences history_geography', 'appliedsci
ences literacy']
```

Normalize numerical features

price

```
In [163]: 1 from sklearn.preprocessing import Normalizer
2 norml=Normalizer()
3 norml.fit(X_train['price'].values.reshape(1,-1))
4 X_train_norml_price=norml.transform(X_train['price'].values.reshape(1,-1))
5 X_test_norml_price=norml.transform(X_test['price'].values.reshape(1,-1))
6 X_cv_norml_price=norml.transform(X_cv['price'].values.reshape(1,-1))
7
8 #X_cv_norm_price=norm.transform(xcv['price'].values.reshape(-1,1))
9 print('X_train_norml_price shape',X_train_norml_price.shape,y_train.shape)
10 print('X_test_norml_price shape',X_test_norml_price.shape,y_test.shape)
11 print('X_cv_norml_price shape',X_cv_norml_price.shape,y_cv.shape)
12
```

```
X_train_norml_price shape (1, 22445) (22445,)
X_test_norml_price shape (1, 16500) (16500,)
X_cv_norml_price shape (1, 11055) (11055,)
```

```
In [ ]: 1
```

teacher_number_of_previously_posted_projects


```
In [164]: 1 norm1.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
2 X_train_norm1_teacher_number_of_previously_posted_projects=norm1.transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
3 X_test_norm1_teacher_number_of_previously_posted_projects=norm1.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
4 X_cv_norm1_teacher_number_of_previously_posted_projects=norm1.transform(X_cv['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
5
6
7 #X_cv_norm_price=norm.transform(xcv['price'].values.reshape(-1,1))
8 print('X_train_norm1_teacher_number_of_previously_posted_projects shape',X_train_norm1_teacher_number_of_previously_posted_projects.shape)
9 print('X_test_norm1_teacher_number_of_previously_posted_projects shape',X_test_norm1_teacher_number_of_previously_posted_projects.shape)
10 #print('X_cv_norm_price shape',xcv_norm_price.shape,ycv.shape)
11 print('X_cv_norm1_teacher_number_of_previously_posted_projects shape',X_cv_norm1_teacher_number_of_previously_posted_projects.shape)
12 #print('X_cv_norm_price shape',xcv_norm_price.shape,ycv.shape)
```

```
X_train_norm1_teacher_number_of_previously_posted_projects shape (1, 22445) (22445,)
X_test_norm1_teacher_number_of_previously_posted_projects shape (1, 16500) (16500,)
X_cv_norm1_teacher_number_of_previously_posted_projects shape (1, 11055) (11055,)
```

SET 1

```
In [165]: 1 # https://stackoverflow.com/a/19710648/4084039
2 # combine all features into one single set
3
4 X_train_norm1_tnoprepst_projects=X_train_norm1_teacher_number_of_previously_posted_projects
5 X_test_norm1_tnoprepst_projects=X_test_norm1_teacher_number_of_previously_posted_projects
6 X_cv_norm1_tnoprepst_projects=X_cv_norm1_teacher_number_of_previously_posted_projects
7
8 X_train_norm1_price= X_train_norm1_price.reshape(22445,1)
9 X_test_norm1_price=X_test_norm1_price.reshape(16500,1)
10 X_cv_norm1_price=X_cv_norm1_price.reshape(11055,1)
11
12 from scipy.sparse import hstack
13 X_tr_set1 = hstack((X_train_essay_bow,X_train_norm1_price,X_train_norm1_tnoprepst_projects))
14 X_tst_set1= hstack((X_test_essay_bow,X_test_norm1_price,X_test_norm1_tnoprepst_projects))
15 X_cr_set1= hstack((X_cv_essay_bow,X_cv_norm1_price,X_cv_norm1_tnoprepst_projects))
16
17
18 print("Final Data matrix")
19 print(X_tr_set1.shape, y_train.shape)
20 print(X_cr_set1.shape, y_cv.shape)
21 print(X_tst_set1.shape, y_test.shape)
22
```

```
Final Data matrix
(22445, 50243) (22445,)
(11055, 50243) (11055,)
(16500, 50243) (16500,)
```

TFIDF vectorization

Eassy

```
In [166]: 1 vect= TfidfVectorizer(min_df=10,ngram_range=(1,4), max_features=50000)
2 vect.fit(X_train['essay'].values)
3 X_train_essay_tfidf = vect.transform(X_train['essay'].values)
4 X_cv_essay_tfidf = vect.transform(X_cv['essay'].values)
5 X_test_essay_tfidf = vect.transform(X_test['essay'].values)
6 essay_tfidf_features=vect.get_feature_names()
7
8 print('X_train_essay_tfidf.shape',X_train_essay_tfidf.shape, y_train.shape)
9 print('X_test_essay_tfidf.shape',X_test_essay_tfidf.shape, y_test.shape)
10 print('X_cv_essay_tfidf.shape',X_cv_essay_tfidf.shape, y_cv.shape)
11 # print('X_cv_essay_bow.shape',X_cv_essay_bow.shape, y_cv.shape)
```

X_train_essay_tfidf.shape (22445, 50000) (22445,)

X_test_essay_tfidf.shape (16500, 50000) (16500,)

X_cv_essay_tfidf.shape (11055, 50000) (11055,)

SET 2

```
In [167]: 1 # https://stackoverflow.com/a/19710648/4084039
2 # combine all features into one single set
3
4 X_train_norm1_tnoprepst_projects=X_train_norm1_teacher_number_of_previously_p
5 X_test_norm1_tnoprepst_projects=X_test_norm1_teacher_number_of_previously_pos
6 X_cv_norm1_tnoprepst_projects=X_cv_norm1_teacher_number_of_previously_posted_
7
8 X_train_norm1_price= X_train_norm1_price.reshape(22445,1)
9 X_test_norm1_price=X_test_norm1_price.reshape(16500,1)
10 X_cv_norm1_price=X_cv_norm1_price.reshape(11055,1)
11
12 from scipy.sparse import hstack
13 X_tr_set2 = hstack((X_train_essay_tfidf,X_train_norm1_price,X_train_norm1_tno
14 X_tst_set2= hstack((X_test_essay_tfidf,X_test_norm1_price,X_test_norm1_tno
15 X_cr_set2= hstack((X_cv_essay_tfidf,X_cv_norm1_price,X_cv_norm1_tnoprepst_pro
16
17
18 print("Final Data matrix")
19 print(X_tr_set2.shape, y_train.shape)
20 print(X_cr_set2.shape, y_cv.shape)
21 print(X_tst_set2.shape, y_test.shape)
22
```

Final Data matrix

(22445, 50243) (22445,)

(11055, 50243) (11055,)

(16500, 50243) (16500,)


1.5 Applying NB on different kind of featurization as mentioned in the instructions

In [228]:

```
1 # please write all the code with proper documentation, and proper titles for
2 # go through documentations and blogs before you start coding
3 # first figure out what to do, and then think about how to do.
4 # reading and understanding error messages will be very much helpfull in debu
5 # when you plot any graph make sure you use
6     # a. Title, that describes your plot, this will be very helpful to the re
7     # b. Legends if needed
8     # c. X-axis label
9     # d. Y-axis label
```

FOR SET 1

```
In [229]: 1 # https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.G
2 from sklearn.model_selection import GridSearchCV
3 from scipy.stats import randint as sp_randint
4 from sklearn.model_selection import RandomizedSearchCV
5 import matplotlib.pyplot as plt
6 from sklearn.naive_bayes import MultinomialNB
7 from sklearn.metrics import roc_auc_score
8 from sklearn import cross_validation
9 import math
10 import warnings
11 warnings.filterwarnings("ignore")
12
13
14 model=MultinomialNB()
15 parameters = {'alpha': sp_randint(0.1,10)}
16 clf = RandomizedSearchCV(estimator = model,param_distributions = parameters,
17 clf.fit(X_tr_set1, y_train)
18
19 results = pd.DataFrame.from_dict(clf.cv_results_)
20 results = results.sort_values(['param_alpha'])
21
22 train_auc= results['mean_train_score']
23 train_auc_std= results['std_train_score']
24 cv_auc = results['mean_test_score']
25 cv_auc_std= results['std_test_score']
26 alpha= results['param_alpha']
27
28 plt.plot(alpha, train_auc, label='Train AUC')
29 # this code is copied from here: https://stackoverflow.com/a/48803361/4084039
30 # plt.gca().fill_between(K, train_auc - train_auc_std,train_auc + train_auc_s
31
32 plt.plot(alpha, cv_auc, label='CV AUC')
33 # this code is copied from here: https://stackoverflow.com/a/48803361/4084039
34 # plt.gca().fill_between(K, cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2
35
36 plt.scatter(alpha, train_auc, label='Train AUC points')
37 plt.scatter(alpha, cv_auc, label='CV AUC points')
38
39
40 plt.legend()
41 plt.xlabel("alpha: hyperparameter")
42 plt.ylabel("AUC")
43 plt.title("Hyper parameter Vs AUC plot")
44 plt.grid()
45 plt.show()
46 results.head()
```



Out[229]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_
7	0.128674	0.000943	0.016334	0.000472	1	{'alpha': 1}	0.7
8	0.139341	0.003682	0.021668	0.003772	1	{'alpha': 1}	0.7
1	0.130674	0.002357	0.019334	0.004714	4	{'alpha': 4}	0.8
3	0.143675	0.006600	0.018334	0.002495	5	{'alpha': 5}	0.8
4	0.159676	0.026839	0.018001	0.002160	6	{'alpha': 6}	0.8

In [230]:

```
1 best_alpha_set1 =clf.best_estimator_  
2 print(best_alpha_set1.alpha)
```


8

FOR SET 2

```

In [231]: 1 model=MultinomialNB()
2 parameters = {'alpha': sp_randint(0.1,10)}
3 clf = RandomizedSearchCV(estimator = model,param_distributions = parameters,
4 clf.fit(X_tr_set2, y_train)
5
6 results = pd.DataFrame.from_dict(clf.cv_results_)
7 results = results.sort_values(['param_alpha'])
8
9 train_auc= results['mean_train_score']
10 train_auc_std= results['std_train_score']
11 cv_auc = results['mean_test_score']
12 cv_auc_std= results['std_test_score']
13 alpha= results['param_alpha']
14
15 plt.plot(alpha, train_auc, label='Train AUC')
16 # this code is copied from here: https://stackoverflow.com/a/48803361/4084039
17 # plt.gca().fill_between(K, train_auc - train_auc_std,train_auc + train_auc_s
18
19 plt.plot(alpha, cv_auc, label='CV AUC')
20 # this code is copied from here: https://stackoverflow.com/a/48803361/4084039
21 # plt.gca().fill_between(K, cv_auc - cv_auc_std,cv_auc + cv_auc_std,alpha=0.2
22
23 plt.scatter(alpha, train_auc, label='Train AUC points')
24 plt.scatter(alpha, cv_auc, label='CV AUC points')
25
26
27 plt.legend()
28 plt.xlabel("alpha: hyperparameter")
29 plt.ylabel("AUC")
30 plt.title("Hyper parameter Vs AUC plot")
31 plt.grid()
32 plt.show()
33 results.head()

```



Out[231]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_test_
2	0.131341	0.001247	0.016001	1.123916e-07	0	{'alpha': 0}	0.8
0	0.133674	0.003300	0.016668	9.427407e-04	4	{'alpha': 4}	0.8
3	0.132008	0.001414	0.016334	4.714266e-04	4	{'alpha': 4}	0.8
9	0.146342	0.010531	0.019001	3.559326e-03	5	{'alpha': 5}	0.8
5	0.130007	0.001633	0.015668	4.714827e-04	6	{'alpha': 6}	0.8

In [232]:

```

1 best_alpha_set2 = clf.best_estimator_
2 # best_alpha_set2.best_alpha
3 print(best_alpha_set2.alpha)

```


4

AUC OF SET 1- BOW

```

In [233]: 1 def batch_predict(clf, data):
2           # roc_auc_score(y_true, y_score) the 2nd parameter should be probability
3           # not the predicted outputs
4
5           y_data_pred = []
6           tr_loop = data.shape[0] - data.shape[0]%1000
7           # consider you X_tr shape is 49041, then your tr_loop will be 49041 - 490
8           # in this for loop we will iterate until the last 1000 multiplier
9           for i in range(0, tr_loop, 1000):
10              y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
11              # we will be predicting for the last data points
12              if data.shape[0]%1000 != 0:
13                  y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])
14
15              return y_data_pred
16
17 clf_set1=MultinomialNB(best_alpha_set1.alpha)
18 clf_set1.fit(X_tr_set1, y_train)
19 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability esti
20 # not the predicted outputs
21
22 y_train_pred = batch_predict(clf_set1, X_tr_set1)
23 y_test_pred = batch_predict(clf_set1, X_tst_set1)
24
25 train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
26 test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
27
28 auc_set1=auc(test_fpr, test_tpr)
29
30 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_t
31 plt.plot(test_fpr, test_tpr, label="train AUC =" +str(auc(test_fpr, test_tpr))
32 plt.legend()
33 plt.xlabel("alpha: hyperparameter")
34 plt.ylabel("AUC")
35 plt.title("ERROR PLOTS")
36 plt.grid()
37 plt.show()

```




AUC OF SET 2- TFIDF

```

In [234]: 1 def batch_predict(clf, data):
2         # roc_auc_score(y_true, y_score) the 2nd parameter should be probability
3         # not the predicted outputs
4
5         y_data_pred = []
6         tr_loop = data.shape[0] - data.shape[0]%1000
7         # consider you X_tr shape is 49041, then your tr_loop will be 49041 - 490
8         # in this for loop we will iterate until the last 1000 multiplier
9         for i in range(0, tr_loop, 1000):
10            y_data_pred.extend(clf.predict_proba(data[i:i+1000])[:,1])
11            # we will be predicting for the last data points
12            if data.shape[0]%1000 != 0:
13                y_data_pred.extend(clf.predict_proba(data[tr_loop:])[:,1])
14
15            return y_data_pred
16
17 clf_set2=MultinomialNB(best_alpha_set2.alpha)
18 clf_set2.fit(X_tr_set1, y_train)
19 # roc_auc_score(y_true, y_score) the 2nd parameter should be probability esti
20 # not the predicted outputs
21
22 y_train_pred = batch_predict(clf_set2, X_tr_set2)
23 y_test_pred = batch_predict(clf_set2, X_tst_set2)
24
25 train_fpr, train_tpr, tr_thresholds = roc_curve(y_train, y_train_pred)
26 test_fpr, test_tpr, te_thresholds = roc_curve(y_test, y_test_pred)
27
28 auc_set2=auc(test_fpr, test_tpr)
29
30 plt.plot(train_fpr, train_tpr, label="train AUC =" +str(auc(train_fpr, train_t
31 plt.plot(test_fpr, test_tpr, label="train AUC =" +str(auc(test_fpr, test_tpr))
32 plt.legend()
33 plt.xlabel("alpha: hyperparameter")
34 plt.ylabel("AUC")
35 plt.title("ERROR PLOTS")
36 plt.grid()
37 plt.show()

```



Confusion matrix

```
In [235]: 1 def find_best_threshold(threshold, fpr, tpr):
2         t = threshold[np.argmax(tpr*(1-fpr))]
3         # (tpr*(1-fpr)) will be maximum if your fpr is very low and tpr is very h
4         print("the maximum value of tpr*(1-fpr)", max(tpr*(1-fpr)), "for threshold
5         return t
6
7 def predict_with_best_t(proba, threshold):
8     predictions = []
9     for i in proba:
10         if i>=threshold:
11             predictions.append(1)
12         else:
13             predictions.append(0)
14     return predictions
15
16 from sklearn.metrics import confusion_matrix
17 best_t = find_best_threshold(tr_thresholds, train_fpr, train_tpr)
18 print("Train confusion matrix")
19 print(confusion_matrix(y_train, predict_with_best_t(y_train_pred, best_t)))
20 print("Test confusion matrix")
21 print(confusion_matrix(y_test, predict_with_best_t(y_test_pred, best_t)))
```

the maximum value of tpr*(1-fpr) 0.5771535030601741 for threshold 0.847

Train confusion matrix

```
[[ 2828   767]
 [ 5020 13830]]
```

Test confusion matrix

```
[[1297 1345]
 [4092 9766]]
```

TOP 20 Features

```
In [236]: 1 import numpy as np
2 feature=[]
3 feature.extend(essay_bow_features)
4 feature.extend(school_state_features)
5 feature.extend(teacher_prefix_features)
6 feature.extend(project_grade_category_features)
7 feature.extend(clean_categories_features)
8 feature.extend(clean_subcategories_features)
9 feature.extend(['price'])
10 feature.extend(['teacher_number_of_previously_posted_projects'])
11
12 print(len(feature))
```

50243

```
In [237]: 1 lst0=np.argsort((clf_set1.feature_log_prob_)[0])[:20]
2 top_20_features_0=np.take(feature,lst0)
3 print(top_20_features_0)
4 print('#####')
5 lst1=np.argsort((clf_set1.feature_log_prob_)[1])[:20]
6 top_20_features_1=np.take(feature,lst1)
7 print(top_20_features_1)
```

['challenge find' 'scooter' 'hokki stools give' 'creative group'
'they represent' 'hokki stools wobble' 'hokki stools would'
'hokki stools would allow' 'scope magazine' 'scientific calculators'
'science unit' 'sugary' 'many students special needs' 'creative active'
'home important' 'science social studies topics' 'allow students listen'
'holder' 'become great readers' 'creative hard']

['literacy_language specialneeds' 'students moderate severe disabilities'
'comfortable these' 'thinking this simply' 'go we need help'
'classroom despite challenges face' 'one effective' 'highest quality'
'thinking this simply cannot' 'functional life' 'supplies donated'
'this simply cannot done' 'as small' 'this simply cannot' 'punishment'
'simply cannot done kids' 'learning sight' 'solve problems make choices'
'problems make choices' 'civics_government literature_writing']

3. Summary

as mentioned in the step 5 of instructions

```
In [238]: 1 ! pip install Prettytable
```

Requirement already satisfied: Prettytable in c:\users\aaa\anaconda3\lib\site-p
ackages (0.7.2)

distributed 1.21.8 requires msgpack, which is not installed.

You are using pip version 10.0.1, however version 20.1 is available.

You should consider upgrading via the 'python -m pip install --upgrade pip' com
mand.

```
In [239]: 1 # http://zetcode.com/python/prettytable/
2 from prettytable import PrettyTable
3 pryt=PrettyTable()
4 pryt.field_names = ["Vectorizer", "Model", "Hyper-Parameter", "AUC"]
5 pryt.add_row([" BOW ", "Naive Bayes", best_alpha_set1.alpha,auc_set1])
6 pryt.add_row([" TFIDF ", "Naive Bayes", best_alpha_set2.alpha, auc_set2])
7 print(pryt)
```

Vectorizer	Model	Hyper-Parameter	AUC
BOW	Naive Bayes	8	0.6620046313811909
TFIDF	Naive Bayes	4	0.6408475705077858

In []:

1
2
3
4
5
6
7