

```
In [42]: 1 import numpy as np
          2 import pandas as pd
          3 from sklearn.datasets import make_classification
          4 import matplotlib.pyplot as plt
```

```
In [43]: 1 X, y = make_classification(n_samples=50000, n_features=15, n_informative=10,
          2                               n_classes=2, weights=[0.7], class_sep=0.7, random
```

```
In [44]: 1 X.shape, y.shape
```

```
Out[44]: ((50000, 15), (50000,))
```

```
In [45]: 1 from sklearn.model_selection import train_test_split
```

```
In [46]: 1 x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25, ra
```

```
In [47]: 1 x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[47]: ((37500, 15), (37500,)), (12500, 15), (12500,))
```

```
In [48]: 1 from sklearn import linear_model
```

```
In [49]: 1 # alpha : float
          2 # Constant that multiplies the regularization term.
          3
          4 # eta0 : double
          5 # The initial learning rate for the 'constant', 'invscaling' or 'adaptive' s
          6
          7 clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', rand
          8 clf
```

```
Out[49]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                        early_stopping=False, epsilon=0.1, eta0=0.0001,
                        fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                        loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                        penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                        tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

In [50]: 1 clf.fit(X=x_train, y=y_train)

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.02 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.03 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.05 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.07 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.08 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.10 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.12 seconds.
-- Epoch 8
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.14 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.17 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.19 seconds.
Convergence after 10 epochs took 0.20 seconds
```

```
Out[50]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0001,
    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

In [51]: 1 clf.coef_, clf.coef_.shape, clf.intercept_

```
Out[51]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867 ,
    0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.18084126,
    0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.02266721]]),
    (1, 15),
    array([ -0.8531383]))
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

Instructions

- Load the datasets(train and test) into the respective arrays
- Initialize the weight_vector and intercept term randomly
- Calculate the initial log loss for the train and test data with the current weight and intercept and store it in a list
- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector
 - Calculate the gradient of the intercept [check this](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing) (<https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing>)
 - Update weights and intercept (check the equation number 32 in the above mentioned pdf (<https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing>)):

$$w^{(t+1)} \leftarrow (1 - \frac{\alpha}{N})w^{(t)} + \alpha x_n(y_n - \sigma((w^{(t)})^T x_n + b^t))$$

$$b^{(t+1)} \leftarrow (b^t + \alpha(y_n - \sigma((w^{(t)})^T x_n + b^t)))$$
 - calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
 - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
 - append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)
- Plot the train and test loss i.e on x-axis the epoch number, and on y-axis the loss
- **GOAL:** compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

```
In [52]: 1 w = np.zeros_like(x_train[0])
          2 b = 0
          3 eta0 = 0.0001
          4 alpha = 0.0001
          5 N = len(x_train)
```

```
In [53]: 1 # write your code to implement SGD as per the above instructions
          2 # please choose the number of iterations on your own
          3 def sigmoid(w,x,b):
          4     return 1/(1+np.exp(-(np.dot(x,w.T)+b)))
```

```
In [54]: 1 # Log loss function Formula  $-(y \log_{10}(y_{\text{hat}}) - (1-y) \log_{10}(1-y_{\text{hat}}))$ 
2 def log_function(w,x,y,b):
3     p = sigmoid(w, x, b)
4     loss =  $-(y * \text{np.log10}(p)) - ((1 - y) * \text{np.log10}(1 - p))$ 
5     return np.mean(loss)
```

In []: 1

```
In [55]: 1 from tqdm import tqdm
2 def trainSGD(x_train, y_train, x_test, y_test, eta0, alpha, epochs, w, b):
3     train_loss=[]
4     test_loss=[]
5     for k in tqdm(range(0, epochs)):
6         for i in range(0,len(x_train)):
7             y = y_train[i]
8             x = x_train[i]
9             w = ((1-eta0*(alpha/N))*w)+((eta0*x)*(y-sigmoid(w,x,b)))
10            b = b+(eta0*(y-sigmoid(w,x,b)))
11
12            # calculating the train loss
13            #     logloss_train=0
14            #     for x in x_train:
15            logloss_train= log_function(w,x_train,y_train,b)
16            train_loss.append(logloss_train)
17
18            #calculating the test loss
19            #     logloss_test=0
20            #     for x in x_train:
21            logloss_test= log_function(w,x_test,y_test,b)
22            test_loss.append(logloss_test)
23
24    return w,b,train_loss, test_loss
```

```
In [56]: 1 epochs=10
          2 w1,b1,trainloss,testloss =trainSGD(x_train, y_train, x_test, y_test,eta0, al
```

```
100%|██████████| 10/10 [00:07<00:00, 1.30it/s]
```

```
In [57]: 1 print(w1)
          2 print(b1)
          3 print(trainloss)
```

```
[-0.42315311  0.19095979 -0.14588118  0.33814991 -0.21196623  0.56525978
-0.44538357 -0.09171679  0.21795314  0.16977398  0.19522044  0.00229554
-0.07781461  0.33882618  0.02214234]
-0.8500967712837224
[0.17546926223702466, 0.16868174436540248, 0.16639953379688374, 0.1653740490192
8135, 0.1648612200408247, 0.16459114506307726, 0.16444479874475637, 0.164364115
2252568, 0.16431912310828212, 0.16429382915597823]
```

In [58]:

```

1 # these are the results we got after we implemented sgd and found the optima
2 w=clf.coef_, b=clf.intercept_

```

Out[58]:

```

(array([[ 0.42336692, -0.18547565,  0.14859036, -0.34144407,  0.2081867 ,
        -0.56016579,  0.45242483,  0.09408813, -0.2092732 , -0.18084126,
        -0.19705191, -0.00421916,  0.0796037 , -0.33852802, -0.02266721]]),
array([0.8531383]))

```

```

1 Difference between sklearn clf coefficient and w1 that is updated by the
  code.

```

In [60]:

```

1 x=w1-clf.coef_
2 y=b1-clf.intercept_
3
4 print(x)
5 print(y)

```

```

[[ 0.0002138  0.00548413  0.00270918 -0.00329416 -0.00377953  0.00509399
   0.00704126  0.00237134  0.00867994 -0.01106728 -0.00183147 -0.00192361
   0.00178909  0.00029817 -0.00052487]]
[0.00304153]

```

In [61]:

```

1 from sklearn.metrics import accuracy_score
2 def pred(w,b, X):
3     N = len(X)
4     predict = []
5     for i in range(N):
6         if sigmoid(w, X[i], b) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(d
7             predict.append(1)
8         else:
9             predict.append(0)
10    return np.array(predict)
11 print(1-np.sum(y_train - pred(w1,b1,x_train))/len(x_train))
12 print(1-np.sum(y_test - pred(w1,b1,x_test))/len(x_test))

```

```

0.95536
0.95296

```

```
In [62]: 1 plt.plot(np.array(trainloss), color='blue', label='Train loss')
2         plt.plot(np.array(testloss), color='red', label='Test loss')
3         plt.ylabel('LogLoss')
4         plt.xlabel('epochs')
5         plt.title('Log loss curve')
6         plt.legend()
7         plt.show()
```

